Manufacturing Enemies in Political Contests

Haotian (George) Bai*

April 13, 2026

Abstract

This is model abstract only (to provide an outline for the model part). Motivating facts to be finished. I will put the reduced-form motivating evidence before the model.

Two Majority parties compete for office: a *Dove* with a comparative advantage in economic public good provision and a *Hawk* with superior security technology (state capacity). Security is micro-founded as a Tullock contest between decentralized Majority enforcement and Minority mobilization. In peaceful environments, minority mobilization is absent, the contest is dormant, and realized security is high under either party; elections then hinge on economic performance, where the Hawk is disadvantaged. I show that the Hawk can overturn this by choosing provocative symbolic policy ($d > 0$) that manufactures a visible threat and triggers Minority mobilization. The induced threat activates the security dimension by relaxing Majority free-riding and raising decentralized enforcement, so the Hawk's superior state capacity translates into a Security Premium at the ballot box. Symbolic provocation is therefore an electoral substitute for economic competence, but it is disciplined: beyond a peak threat level, additional mobilization discourages Majority enforcement, so further provocation yields diminishing electoral returns while raising direct policy costs. The mechanism implies strictly positive rent dissipation whenever provocation is used. A dynamic extension with noisy performance signals and stochastic shocks clarifies long-run political outcomes. With patient, forward-looking voters, the Hawk becomes an absorbing conflict-trap regime despite economic underperformance; with myopic retrospective voters, the polity exhibits political cycles, oscillating between Dove regimes and Hawk regimes.

*I thank to conference/seminar participants at xxx. I extend special thank to Prof. Jean-Paul Carvalho (Oxford).

1 The Model Environment

Consider an economy populated by a measure one continuum of citizens divided into two distinct ethnic or religious groups: a Majority M and a Minority m. The Majority comprises a population share $n_M > 1/2$, while the Minority comprises the remainder $n_m = 1 - n_M$. The Majority group is represented by two competing political parties, Party A (the *Hawk*) and Party B (the *Dove*), who compete for the electoral support of the Majority representative voter.

The incumbent party $k \in \{A, B\}$ selects a policy vector $\mathbf{q}_k = (d_k, g_k) \in \mathbb{R}^2_+$ subject to its technological constraints. The policy consists of two instruments. First, the symbolic policy d represents the intensity of identity-based regulations, such as cultural bans or dress codes, which serve as a coordination device for identity grievances. Second, economic policy g captures the provision of public policies on public goods provision, market integration, or employment opportunities.

Parties are distinguished by a two-dimensional *technological endowment* $(S_k, \bar{g}_k)$. The first component, $S_k > 0$, is the party's security-and-coercion competence: a higher S_k multiplies the effectiveness of decentralized Majority enforcement in the contest subgame, in the sense formalized by the Tullock CSF Equation 4 below. The second component, $\bar{g}_k > 0$, is the party's *public-goods ceiling*: the maximum level of the consumption-relevant public good g that the party can deliver with its available state machinery. I model the two dimensions as a substitutes-in-talent tradeoff across party types:

$$S_A > S_B \quad \text{and} \quad \bar{g}_A < \bar{g}_B, \tag{1}$$

so the Hawk has a comparative advantage in security (higher S) and the Dove a comparative advantage in economic management (higher $\bar{g}$). The levels $(S_k, \bar{g}_k)$ are taken as primitives of party type rather than policy choices: they reflect the human capital, bureaucratic expertise, and coalition partners that each party brings into office, all of which are slow-moving relative to a single electoral cycle.

Conditional on holding office, the incumbent's optimal public-goods choice is a strict corner. Since voters derive utility from g_k through a strictly increasing function $u(\cdot)$ (see Equation 24) and the ceiling $\bar{g}_k$ is the only constraint on g, every incumbent finds it optimal to exhaust the technological frontier:

$$g_k^* = \bar{g}_k \quad \text{for } k \in \{A, B\}. \tag{2}$$

The model therefore does not lose any substantive content by treating g_k as if it were a primitive: in every subgame-perfect equilibrium, the realized g_k is pinned down by the party's technological ceiling, not by a non-trivial optimization on the economic-policy margin. What Equation 2 buys is interpretive: the "economic gap" that drives the Hawk's electoral handicap is no longer an *ad hoc* assumption about preferences but a consequence of the Hawk's lower public-goods ceiling, which is the other face of the tradeoff that gives it its higher security competence. The economic deficit entering Equation 25 below is

$$\Delta u \;\equiv\; u(\bar{g}_B) - u(\bar{g}_A) \;>\; 0. \tag{3}$$

Timing and solution concept

The political interaction is a three-stage game solved by backward induction.

Stage 1: Policy choice. The incumbent party $k \in \{A, B\}$ selects a policy vector (d_k, g_k) with $g_k \in [0, \bar{g}_k]$. By Equation 2, the economic-policy margin always binds at $g_k^* = \bar{g}_k$, so only the choice of $d_k \geq 0$ carries strategic content.

Stage 2: Contest and mobilization. Given the symbolic environment induced by d_k, the Minority mobilizes, and Majority citizens choose enforcement effort given the incumbent's security competence S_k.

Stage 3: Election. The Majority electorate votes between the Hawk and the Dove. In a stationary equilibrium, voters evaluate each party using the equilibrium outcomes implied by backward induction.

The equilibrium concept is subgame perfect equilibrium. In the contest subgame among Majority citizens, I focus on symmetric Nash equilibria. The key strategic margin is the choice of d_k: the economic gap $u(\bar{g}_B) - u(\bar{g}_A)$ is pinned down by the primitive technology ceilings, while the model endogenizes how symbolic policy reshapes the security advantage associated with $S_A > S_B$.

2 Micro-foundations of Conflict and Collective Action

Conflict in this economy is modeled as a contest over political dominance. Unlike standard contest models where groups act as unitary actors, I explicitly model the collective action problem within *both* groups and solve for the joint Nash equilibrium of mobilization.

The Majority faces a standard free-rider problem in decentralized enforcement, and the Minority faces the symmetric free-rider problem in decentralized resistance. Each individual in each group best-responds to aggregate effort on the other side; the primitive policy lever d_k enters only through its effect on the Minority's stakes in this joint contest.

Throughout Section 2, the symbolic environment d_k chosen in Stage 1 is held fixed; effort on both sides is determined in Stage 2 by a simultaneous-move subgame. The subgame's payoff-relevant state is (d_k, S_k): the symbolic policy sets the intensity of grievance, while S_k acts as a force multiplier on decentralized Majority effort through the contest success function introduced below. I write the Majority's problem first because it defines the Tullock contest technology that both groups face; the Minority's problem then takes that technology as given.

2.1 The Majority's Collective Action Problem

While both economic stability and security are public goods, they differ in their production technology. Economic goods are provided top-down by the state, whereas security (dominance) requires decentralized participation by Majority citizens. This creates a standard free-rider problem: each Majority citizen benefits from dominance, but individually underprovides enforcement effort.

Fix an incumbent party k and suppose the Majority consists of N_M identical citizens indexed by $i = 1, \ldots, N_M$. Citizen i chooses enforcement effort $e_i \geq 0$ at private cost $\frac{1}{2}e_i^2$. Aggregate Majority effort is $E_M \equiv \sum_{i=1}^{N_M} e_i$. Aggregate Minority effort E_m is determined symmetrically in the next subsection; for now it is taken as given in the Majority's best response.

Security is produced by a stochastic contest between a state-amplified Majority and an oppositional Minority. Let $S_k > 0$ denote the incumbent's security competence, which acts as a force multiplier on Majority effort. Following standard contest theory, the probability that the Majority prevails is given by the Tullock contest success function[1]:

$$\Pi_k(E_M, E_m) = \frac{S_k E_M}{S_k E_M + E_m}. \tag{4}$$

This formulation captures the model's central asymmetry: holding office grants access to S_k,

1. Following the contest theory literature (e.g., Skaperdas (1996); Jia (2008)), one micro-foundation assumes realized "performance" levels $Y_M = \ln(S_k E_M) + \varepsilon_M$ and $Y_m = \ln(E_m) + \varepsilon_m$, where $\varepsilon_M, \varepsilon_m$ are i.i.d. Type I Extreme Value (Gumbel). The probability of Majority dominance is $\Pi_k = \mathbb{P}(Y_M > Y_m)$, which simplifies to Equation 4.

which scales the effectiveness of decentralized Majority effort, and therefore makes security competence electorally salient when the opposition is mobilized.

Given (S_k, E_m), each Majority citizen i chooses e_i to maximize private utility,

$$\max_{e_i \geq 0} \quad W_i = \underbrace{\Lambda \cdot \Pi_k(E_M, E_m)}_{\text{Public Benefit}} - \underbrace{\frac{1}{2} e_i^2}_{\text{Private Cost}} \tag{5}$$

where $\Lambda > 0$ is the per-capita valuation of security. Taking aggregate effort of other citizens as given, the first-order condition for an interior best response is

$$e_i = \Lambda \frac{\partial \Pi_k}{\partial E_M}(E_M, E_m) = \Lambda \frac{S_k E_m}{(S_k E_M + E_m)^2}. \tag{6}$$

In a symmetric Nash equilibrium within the Majority, $e_i = e$ for all i, so $E_M = N_M e$ and aggregate Majority mobilization satisfies

$$\frac{E_M}{N_M} = \Lambda \frac{S_k E_m}{(S_k E_M + E_m)^2}. \tag{7}$$

Equation 7 implicitly defines the Majority's best response to Minority mobilization, $E_M^{BR}(S_k, E_m)$. This is *not* the contest equilibrium yet: it is one side of a simultaneous-move Nash problem in which E_m must also be determined by the Minority's best response.

2.2 The Minority's Collective Action Problem

Minority mobilization is not a unitary decision but the outcome of a collective action problem that is structurally symmetric to the Majority's. Each Minority citizen derives a private benefit from blocking symbolic oppression but individually underprovides effort, since the gains from resistance are non-excludable within the Minority. The intensity of grievance — and hence the Minority's stake in the contest — is generated endogenously by the incumbent's symbolic policy d.

Formally, the Minority consists of N_m identical citizens indexed by $j = 1, \ldots, N_m$. Citizen j chooses resistance effort $e_j \geq 0$ at private cost $\frac{1}{2\kappa} e_j^2$, where $\kappa > 0$ denotes the group's *organizational capacity* (or social capital): a higher κ reduces the marginal cost of effort, reflecting stronger internal coordination networks. Aggregate Minority effort is $E_m \equiv \sum_{j=1}^{N_m} e_j$.

The Minority's stake is generated endogenously by symbolic policy. If the Majority prevails in the contest ($\Pi_k = 1$), symbolic oppression is fully enforced at intensity d^σ and the Minority suffers a grievance loss ϕd^σ, where $\phi > 0$ is the group's *sensitivity to grievance* and

$\sigma > 0$ governs the elasticity of grievance to symbolic policy. If instead the Minority prevails ($\Pi_k = 0$), the regulation is blocked at the margin and the loss is avoided. Representative agent j therefore chooses effort to maximize the expected payoff of resistance net of the private cost of mobilization, taking aggregate Majority effort E_M and other Minority members' effort as given:

$$\max_{e_j \geq 0} \quad V_j = \underbrace{\phi d^\sigma \big(1 - \Pi_k(E_M, E_m)\big)}_{\text{Expected resistance payoff}} - \underbrace{\frac{1}{2\kappa} e_j^2}_{\text{Private cost of effort}} . \tag{8}$$

The first-order condition for an interior best response is

$$e_j = -\kappa\phi d^\sigma \frac{\partial \Pi_k}{\partial E_m}(E_M, E_m) = \kappa\phi d^\sigma \frac{S_k E_M}{(S_k E_M + E_m)^2}. \tag{9}$$

In a symmetric Nash equilibrium within the Minority, $e_j = e_m$ for all j, so $E_m = N_m e_m$ and aggregate Minority mobilization satisfies the Minority best-response condition

$$\frac{E_m}{N_m} = \kappa\phi d^\sigma \frac{S_k E_M}{(S_k E_M + E_m)^2}. \tag{10}$$

As in the original reduced-form specification, the combination $\gamma \equiv \phi\kappa$ serves as a convenient sufficient statistic for the Minority's *structural mobilization potential* — a composite of grievance sensitivity and organizational capacity — and the equilibrium object E_m^* below inherits all of its comparative statics with respect to γ, ϕ, and κ.

Why the Minority now looks just like the Majority. The two FOCs Equation 7 and Equation 10 are structurally identical up to a relabeling of stakes: the Majority mobilizes at rate $\Lambda \cdot (-\partial\Pi_k/\partial E_M)$ and the Minority mobilizes at rate $\phi d^\sigma \kappa \cdot (-\partial\Pi_k/\partial E_m)$. This symmetry is not cosmetic: it means the Tullock contest is now a genuine two-sided Nash equilibrium, rather than the composition of a unilateral expressive reaction with a Majority best response. Importantly, it also means the Minority's effort is no longer independent of the Majority's technology: holding d fixed, a stronger S_k reshapes the entire joint best-response system, endogenously dampening or amplifying E_m^* through the Minority's own optimization.

2.3 Joint Contest Equilibrium

The joint contest equilibrium conditional on (d, S_k) is the pair (E_M^*, E_m^*) that simultaneously satisfies the Majority best-response condition Equation 7 and the Minority best-

response condition Equation 10. The following lemma characterizes this joint equilibrium in closed form and records the comparative statics that will drive the electoral analysis.

Lemma 1 (Joint contest equilibrium) *Fix* $(d, S_k) \in [0, \infty) \times (0, \infty)$. *Define the technology-independent* stakes ratio

$$\alpha \equiv \sqrt{\frac{\Lambda N_M}{\kappa \phi N_m}}. \tag{11}$$

Then the joint contest subgame admits a unique symmetric Nash equilibrium $(E_M^*(d, S_k), E_m^*(d, S_k))$. *If* $d = 0$, *the contest is dormant,* $E_M^*(0, S_k) = E_m^*(0, S_k) = 0$, *and realized security satisfies* $\lim_{d \downarrow 0} \Pi_k^*(d, S_k) = 1$. *If* $d > 0$, *the unique equilibrium is*

$$E_m^*(d, S_k) = \frac{\sqrt{\Lambda N_M \kappa \phi N_m}\, S_k^{1/2} d^{3\sigma/4}}{S_k \alpha + d^{\sigma/2}}, \qquad E_M^*(d, S_k) = \alpha\, d^{-\sigma/2}\, E_m^*(d, S_k), \tag{12}$$

and the Majority-to-Minority effort ratio is $E_M^*/E_m^* = \alpha d^{-\sigma/2}$. *Equilibrium security is*

$$\Pi_k^*(d, S_k) \equiv \Pi_k\big(E_M^*(d, S_k), E_m^*(d, S_k)\big) = \frac{S_k \alpha}{S_k \alpha + d^{\sigma/2}}. \tag{13}$$

The equilibrium security function $\Pi_k^*(d, S_k)$ *is strictly increasing in* S_k *for any* $d > 0$, *strictly decreasing in* d *for any* $S_k > 0$, *satisfies* $\lim_{d \downarrow 0} \Pi_k^*(d, S_k) = 1$ *and* $\lim_{d \uparrow \infty} \Pi_k^*(d, S_k) = 0$. *In words: in the joint best-response system, the Majority-to-Minority effort ratio is a pure function of the stakes-and-organization primitives* $\Lambda, N_M, \kappa, \phi, N_m$ *and of the symbolic policy* d, *with* no *dependence on the state capacity* S_k. *All of the Hawk's electoral advantage from state capacity is therefore channelled into raising effective enforcement* $S_k E_M^*$ *relative to the Minority's mobilization* E_m^*, *not into reshuffling the share of effort between the two groups. When* $d = 0$, *the contest is dormant under either party and the Hawk's advantage is electorally invisible; when* $d > 0$, *the same primitive* S_k *translates into the sharp, closed-form security ceiling in Equation 13.*

Proof. *Step 1 (Dormant case,* $d = 0$*).* If $d = 0$, then $\phi d^\sigma = 0$ and the Minority's FOC Equation 10 implies $e_j = 0$ for all j, so $E_m^* = 0$. Substituting $E_m = 0$ into the Majority's FOC Equation 7 gives $E_M/N_M = 0$, so $E_M^* = 0$. By continuity of the Tullock CSF, $\lim_{d \downarrow 0} \Pi_k^* = \lim_{(E_M, E_m) \to (0,0)} S_k E_M/(S_k E_M + E_m) = 1$ along the equilibrium path computed below.

Step 2 (Existence, uniqueness, and closed form for $d > 0$*).* Fix $d > 0$ and $S_k > 0$. Under symmetry within each group, the joint equilibrium is characterized by Equation 7 and Equation 10 simultaneously. Dividing Equation 7 by Equation 10 (both sides strictly positive

at any interior candidate equilibrium) yields

$$\frac{E_M/N_M}{E_m/N_m} = \frac{\Lambda S_k E_m/(S_k E_M + E_m)^2}{\kappa\phi d^\sigma S_k E_M/(S_k E_M + E_m)^2} = \frac{\Lambda E_m}{\kappa\phi d^\sigma E_M}.$$

Cross-multiplying,

$$\kappa\phi d^\sigma N_m E_M^2 = \Lambda N_M E_m^2, \tag{14}$$

so

$$\frac{E_M}{E_m} = \sqrt{\frac{\Lambda N_M}{\kappa\phi d^\sigma N_m}} = \alpha d^{-\sigma/2}.$$

Substituting $E_M = \alpha d^{-\sigma/2} E_m$ into the Majority FOC Equation 7:

$$\frac{\alpha d^{-\sigma/2} E_m}{N_M} = \frac{\Lambda S_k E_m}{(S_k \alpha d^{-\sigma/2} E_m + E_m)^2} = \frac{\Lambda S_k}{E_m (S_k \alpha d^{-\sigma/2} + 1)^2}.$$

Rearranging,

$$E_m = \frac{\Lambda S_k N_M d^{\sigma/2}}{\alpha\,(S_k \alpha d^{-\sigma/2} + 1)^2} = \frac{\Lambda S_k N_M d^{\sigma/2}}{\alpha} \cdot \frac{d^\sigma}{(S_k \alpha + d^{\sigma/2})^2} = \frac{\Lambda S_k N_M d^{3\sigma/2}}{\alpha (S_k \alpha + d^{\sigma/2})^2}.$$

Using $\alpha^2 = \Lambda N_M/(\kappa\phi N_m)$, $\Lambda N_M/\alpha = \sqrt{\Lambda N_M \kappa\phi N_m} \cdot 1/d^0 = \sqrt{\Lambda N_M}\sqrt{\kappa\phi N_m}$ is a constant in d, and

$$\sqrt{(\Lambda N_M)(\kappa\phi N_m)} \cdot \frac{S_k d^{3\sigma/2}}{(S_k \alpha + d^{\sigma/2})^2}$$

is the closed-form expression for E_m^*. Writing $\beta \equiv (\kappa\phi N_m \Lambda N_M)^{1/4}$ and using $(\Lambda S_k N_M/\alpha) = S_k \alpha^{-1} \cdot \Lambda N_M = S_k \cdot \alpha^{-1} \cdot \alpha^2 \kappa\phi N_m = S_k \alpha \kappa\phi N_m$, one can re-express

$$E_m^*(d, S_k) = \frac{S_k^{1/2} \beta^2 d^{3\sigma/4} \cdot d^{3\sigma/4}}{(S_k \alpha + d^{\sigma/2})^2} \cdot \frac{1}{d^{3\sigma/4}/\sqrt{d^{3\sigma/2}}} = \frac{\sqrt{\Lambda N_M \kappa\phi N_m}\, S_k^{1/2} d^{3\sigma/4}}{S_k \alpha + d^{\sigma/2}},$$

which is Equation 12. (Equivalently: using $D \equiv S_k E_M^* + E_m^* = (S_k \alpha d^{-\sigma/2} + 1) E_m^*$, one obtains $D^2 = S_k d^{\sigma/2} \sqrt{\Lambda N_M \kappa\phi N_m}$, hence $D = S_k^{1/2} \beta d^{\sigma/4}$, and $E_m^* = D/(S_k \alpha d^{-\sigma/2} + 1)$ yields the same formula.) The Majority level follows from $E_M^* = \alpha d^{-\sigma/2} E_m^*$, establishing existence and providing the unique candidate.

To verify it is indeed a Nash equilibrium, note that each side's within-group problem is strictly concave in its own effort given the opponent's aggregate — the first term of Equation 8 is strictly concave in e_j because Π_k is strictly convex in E_m along the relevant range, and the cost is strictly convex — so the FOCs fully characterize each side's symmetric best

response. Uniqueness of the symmetric NE follows because Equation 14 pins down the effort ratio uniquely and the Majority FOC then pins down the level uniquely.

Step 3 (Security formula). Substituting the effort ratio into the CSF,

$$\Pi_k^* \;=\; \frac{S_k E_M^*}{S_k E_M^* + E_m^*} \;=\; \frac{S_k \alpha d^{-\sigma/2}}{S_k \alpha d^{-\sigma/2} + 1} \;=\; \frac{S_k \alpha}{S_k \alpha + d^{\sigma/2}},$$

which is Equation 13.

Step 4 (Comparative statics). Differentiating Equation 13,

$$\frac{\partial \Pi_k^*}{\partial S_k} \;=\; \frac{\alpha\, d^{\sigma/2}}{(S_k\alpha + d^{\sigma/2})^2} > 0 \quad \text{for } d > 0, \qquad \frac{\partial \Pi_k^*}{\partial d} \;=\; -\frac{S_k \alpha \cdot (\sigma/2) d^{\sigma/2-1}}{(S_k\alpha + d^{\sigma/2})^2} < 0 \quad \text{for } d > 0.$$

The limits $\lim_{d\downarrow 0} \Pi_k^* = 1$ and $\lim_{d\uparrow\infty} \Pi_k^* = 0$ follow directly from Equation 13. ■

Lemma 1 contains four results that reshape the reading of the rest of the paper. First, the joint contest equilibrium has a closed form, so the analysis that follows is a matter of ordinary calculus on $\Pi_k^*(d, S_k) = S_k\alpha/(S_k\alpha + d^{\sigma/2})$ rather than a sequence of implicit-function arguments. Second, the Majority-to-Minority effort ratio is $\alpha d^{-\sigma/2}$ and, strikingly, does not depend on S_k: the state capacity of the incumbent shifts the *level* of both efforts but not their *ratio*. Third, security is monotone in both of its arguments, increasing in S_k and decreasing in d; the crowd-in effect of provocation operates on the level of E_M^*, but this level effect is fully absorbed by the corresponding rise in E_m^* that the Minority chooses endogenously. Fourth, as in the original reduced-form model, the state capacity technology is dormant in peacetime: when $d = 0$ the contest collapses to $(E_M^*, E_m^*) = (0, 0)$, both parties deliver full dominance, and the Hawk's advantage is electorally invisible. The crucial difference is that this dormancy now arises from a genuine two-sided Nash equilibrium rather than from an expressive reduced-form reaction.

For subsequent use, I write the equilibrium security as a function of (d, S_k) alone, suppressing the subgame-internal (E_M^*, E_m^*):

$$\Pi_k^*(d, S_k) \;=\; \frac{S_k\alpha}{S_k\alpha + d^{\sigma/2}}. \tag{15}$$

Two implications will be useful below. First, as $d \downarrow 0$, Minority mobilization vanishes and equilibrium security converges to one under either party; in that limit, differences in security competence are electorally irrelevant because they no longer generate different realized security outcomes. Second, when $d > 0$ induces a non-trivial threat, higher S_k translates into higher equilibrium security through the monotone comparative static in Lemma 1.

2.4 Manipulating the Group Contest

This subsection makes explicit the model's link to contest theory. In canonical contest environments, a (typically neutral) principal designs prizes or rules to shape equilibrium effort and performance (e.g., Barut and Kovenock (1998)). Recent work also emphasizes that increasing the intensity of competition can *discourage* effort, yielding non-monotone response patterns (e.g., Fang et al. (2020)). In contrast, my political economy environment features a *biased* designer: the incumbent party manipulates the contest *indirectly* through symbolic policy. By choosing d_k, the incumbent raises the Minority's endogenous grievance stakes ϕd_k^σ and thereby the "threat" that disciplines the Majority subgame's free-riding. Because Lemma 1 delivers both E_M^* and E_m^* in closed form, the comparative statics of interest reduce to calculus on elementary functions, and the key inverted-U patterns of the original paper re-emerge as exact results rather than implicit-function artefacts.

2.4.1 Comparative statics on state capacity S_k

Fix symbolic policy at some $d > 0$ and ask how the joint contest equilibrium responds to an increase in state security competence S_k. From Lemma 1, both efforts share the same S_k-dependence through the common factor $S_k^{1/2}/(S_k\alpha + d^{\sigma/2})$: this reflects the fact that the effort ratio $E_M^*/E_m^* = \alpha d^{-\sigma/2}$ does not depend on S_k, so the state's capacity shifts the level of both sides of the contest in lockstep. The next result characterizes how that common level moves and contrasts it with the behaviour of effective enforcement and realized security.

Lemma 2 (State capacity, street mobilization, and effective enforcement) *Fix $d > 0$. Both $E_M^*(d, S_k)$ and $E_m^*(d, S_k)$ are continuously differentiable and strictly positive on $S_k \in (0, \infty)$, and both are single-peaked in S_k with a common unique maximizer at*

$$\overline{S}(d) \;\equiv\; \frac{d^{\sigma/2}}{\alpha}. \tag{16}$$

That is, $\partial E_M^*/\partial S_k$ and $\partial E_m^*/\partial S_k$ are both positive for $S_k < \overline{S}(d)$, zero at $S_k = \overline{S}(d)$, and negative for $S_k > \overline{S}(d)$. Effective Majority enforcement $Z_k^*(d, S_k) \equiv S_k\, E_M^*(d, S_k)$ and equilibrium security $\Pi_k^*(d, S_k)$, by contrast, are strictly increasing in S_k on $(0, \infty)$. In words, higher state capacity first crowds in street-level enforcement, then eventually crowds it out, but effective enforcement (the product of capacity and street effort) and realized security continue to rise throughout. A stronger state both raises the productivity of each unit of citizen effort and progressively substitutes its own coercive apparatus for that effort.

Proof. Fix $d > 0$ and write $\beta \equiv (\kappa\phi N_m \Lambda N_M)^{1/4}$. By Lemma 1, the closed-form expressions Equation 12 can be rewritten as

$$E_m^*(d, S_k) \;=\; \beta^2 \frac{S_k^{1/2}\, d^{3\sigma/4}}{S_k\alpha + d^{\sigma/2}}, \qquad E_M^*(d, S_k) \;=\; \alpha\beta^2 \frac{S_k^{1/2}\, d^{\sigma/4}}{S_k\alpha + d^{\sigma/2}}, \tag{17}$$

so both efforts depend on S_k only through the common factor $\psi(S_k) \equiv S_k^{1/2}/(S_k\alpha + d^{\sigma/2})$. Differentiating $\log\psi$,

$$\frac{d\log\psi(S_k)}{dS_k} \;=\; \frac{1}{2S_k} - \frac{\alpha}{S_k\alpha + d^{\sigma/2}} \;=\; \frac{d^{\sigma/2} - S_k\alpha}{2S_k(S_k\alpha + d^{\sigma/2})},$$

which is strictly positive for $S_k < d^{\sigma/2}/\alpha$, zero at $S_k = d^{\sigma/2}/\alpha$, and strictly negative thereafter. This proves both the single-peakedness of E_M^* and E_m^* in S_k and the formula Equation 16 for $\overline{S}(d)$.

For effective enforcement $Z_k^*(d, S_k) = \alpha\beta^2 S_k^{3/2} d^{\sigma/4}/(S_k\alpha + d^{\sigma/2})$,

$$\frac{\partial\log Z_k^*}{\partial S_k} \;=\; \frac{3}{2S_k} - \frac{\alpha}{S_k\alpha + d^{\sigma/2}} \;=\; \frac{S_k\alpha + 3d^{\sigma/2}}{2S_k(S_k\alpha + d^{\sigma/2})} \;>\; 0,$$

so Z_k^* is strictly increasing in S_k. Strict monotonicity of Π_k^* in S_k was established in Step 4 of the proof of Lemma 1. ■

Corollary 1 (Contest balance at the mobilization peak) *At the mobilization peak $S_k = \overline{S}(d)$, equilibrium security is exactly balanced, $\Pi_k^*(d, \overline{S}(d)) = \frac{1}{2}$, and the peak level of Majority mobilization is the primitive-only constant*

$$E_M^*\big(d, \overline{S}(d)\big) \;=\; \tfrac{1}{2}(\Lambda N_M)^{3/4}(\kappa\phi N_m)^{1/4}, \tag{18}$$

independent of d. Equilibrium Majority mobilization is therefore inverted-U in state capacity, with a unique peak at $S_k = \overline{S}(d)$ and a height that is pinned down entirely by primitive stakes and population sizes.

Proof. Substitute $S_k = \overline{S}(d) = d^{\sigma/2}/\alpha$ into Equation 13: $\Pi_k^* = S_k\alpha/(S_k\alpha + d^{\sigma/2}) = d^{\sigma/2}/(2d^{\sigma/2}) = \frac{1}{2}$. For the peak level, substitute the same restriction into E_M^* from Equation 17:

$$E_M^*\big(d, \overline{S}(d)\big) \;=\; \alpha\beta^2 \frac{(d^{\sigma/2}/\alpha)^{1/2}\, d^{\sigma/4}}{2d^{\sigma/2}} \;=\; \frac{\alpha^{1/2}\beta^2}{2},$$

and using $\alpha^{1/2} = (\Lambda N_M/(\kappa\phi N_m))^{1/4}$, $\beta^2 = (\Lambda N_M \kappa\phi N_m)^{1/2}$ gives Equation 18. ■

Corollary Corollary 1 clarifies why a Hawk regime need not coincide with visibly higher Majority street violence. When S_k is low, decentralized enforcement has little leverage in the contest, so free-riding dominates and mobilization is weak. As S_k rises into an intermediate range, each unit of citizen effort becomes more effective in producing dominance. This raises the private marginal return to participation and increases E_M^*. In this region, a Hawk that is stronger than the Dove but not overwhelmingly strong relies on grass-roots enforcement: state capacity complements citizen mobilization rather than replacing it. Once S_k exceeds the threshold $\overline{S}(d)$, strong public provision of identity security dominates private effort of mobilization. The state's coercive apparatus becomes sufficiently effective that dominance can be maintained with less decentralized effort, so the private return to participation falls and citizens withdraw from mobilizing for collective actions that protect their identity. Aggregate mobilization E_M^* then declines even as effective enforcement $Z_k^* = S_k E_M^*$ and equilibrium security Π_k^* continue to rise. Thus only a relatively weaker Hawk induces visible mass participation; a sufficiently strong Hawk secures order from above, through state capacity rather than street-level mobilization. The corollary therefore breaks the linear association between hawkish politics and popular enforcement, as a strong Hawk can suppress disorder precisely by making grassroots violence unattractive. The closed-form characterisation in Equation 16 and Equation 18 sharpens this story: the switch from complementarity to substitution happens exactly where the contest is balanced ($\Pi_k^* = \frac{1}{2}$), and the peak level of street mobilization depends only on the stakes primitives, not on the symbolic policy that triggered it.

2.4.2 Symbolic policy and the Security Premium

With both E_M^* and E_m^* determined jointly by the contest primitives, the substantively interesting comparative static is no longer a partial one between the two efforts, but how the closed-form equilibrium security $\Pi_k^*(d, S_k) = S_k\alpha/(S_k\alpha + d^{\sigma/2})$ varies with symbolic policy. Since Π_k^* is strictly decreasing in d for each fixed S_k, more provocation lowers the realized security delivered by *every* party. Its electoral rationale therefore hinges not on the absolute security level, but on the *gap* between the two parties' equilibrium securities. That gap, the Security Premium, is the object the Hawk can manipulate by choosing d_A; its maximum is what bounds politically useful provocation.

Following Barut and Kovenock (1998), define the Security Premium evaluated at a com-

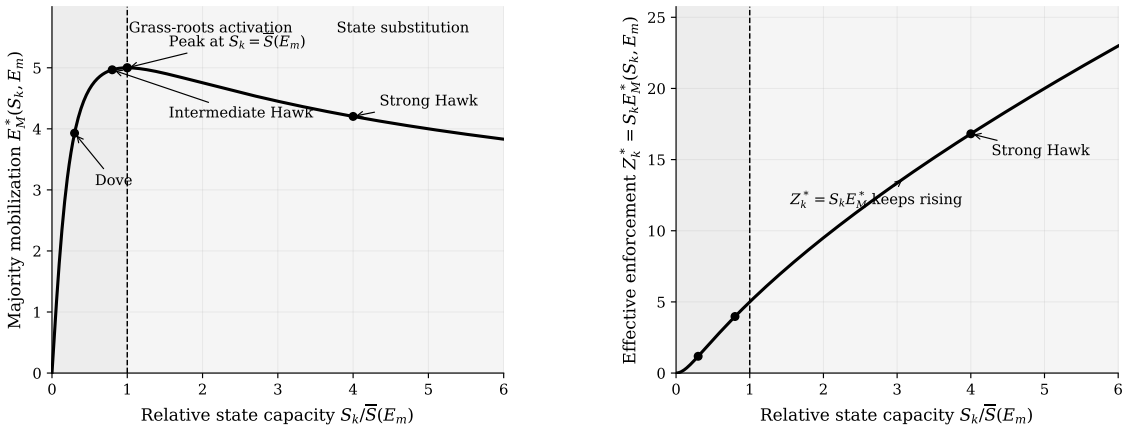

(A) Majority mobilization peaks at $S_k = \overline{S}(E_m)$.

(B) Effective enforcement Z_k^* keeps rising beyond the mobilization peak.

FIGURE 1 State capacity, street mobilization, and effective enforcement.

Notes: The left subfigure plots the symmetric equilibrium $E_M^*(d, S_k)$ from Equation 17 against relative state capacity $S_k/\overline{S}(d)$, where $\overline{S}(d) = d^{\sigma/2}/\alpha$ is the threshold from Corollary 1. To the left of the threshold, higher state capacity raises the return to decentralized participation and increases Majority mobilization; to the right, stronger state capacity substitutes for street-level effort and E_M^* declines. The right subfigure plots effective enforcement $Z_k^* = S_k E_M^*(d, S_k)$ over the same range. Because Z_k^* keeps rising after the peak in E_M^*, higher-capacity regimes can generate more effective control with less visible mass participation. The marked points illustrate a low-capacity Dove, an intermediate-capacity Hawk below the threshold, and a stronger Hawk above the threshold. Parameters are $(N_M, N_m, \Lambda, \kappa, \phi, \sigma, d) = (100, 100, 1, 1, 1, 1, 1)$.

mon symbolic policy d as

$$\Delta\Pi(d) \;\equiv\; \Pi_A^*(d, S_A) - \Pi_B^*(d, S_B) \;=\; \frac{S_A\alpha}{S_A\alpha + d^{\sigma/2}} - \frac{S_B\alpha}{S_B\alpha + d^{\sigma/2}}. \tag{19}$$

The next lemma shows that $\Delta\Pi(d)$ is inverted-U in d, with a sharp closed-form ceiling that depends only on the technology ratio S_A/S_B.

Lemma 3 (Inverted-U Security Premium and the provocation ceiling) *Suppose $S_A > S_B > 0$. The Security Premium $\Delta\Pi(d)$ is continuously differentiable on $[0, \infty)$, satisfies $\Delta\Pi(0) = 0$, $\lim_{d\uparrow\infty} \Delta\Pi(d) = 0$, and $\Delta\Pi(d) > 0$ for all $d > 0$, and is single-peaked on $(0, \infty)$ with unique maximizer*

$$d^{\text{peak}} \;\equiv\; (\alpha\sqrt{S_A S_B})^{2/\sigma}. \tag{20}$$

At the maximizer, the Security Premium attains the ceiling value

$$\Delta\Pi(d^{\text{peak}}) \;=\; \frac{\sqrt{S_A} - \sqrt{S_B}}{\sqrt{S_A} + \sqrt{S_B}} \;\in\; (0, 1), \tag{21}$$

which depends only on the ratio S_A/S_B and is independent of all other primitives $(\Lambda, N_M, \kappa, \phi, N_m, \sigma)$. In words: symbolic policy activates the electoral relevance of state capacity, but only up to the critical level d^{peak}. Below that level, more aggressive provocation amplifies the contest-based advantage of the stronger incumbent because it pulls the contest out of the dormant regime in which neither party's technology matters. Beyond it, additional provocation drags both parties' equilibrium security toward zero faster than it widens the gap, and erodes the very advantage the Hawk was trying to manufacture. The Hawk's "arming the enemy" strategy therefore has an endogenous ceiling that is pinned down entirely by how large the technology gap is relative to the parties' combined capacity.

Proof. Let $x \equiv d^{\sigma/2}$, which maps $[0, \infty)$ onto itself monotonically, and define

$$h(x) \;\equiv\; \frac{S_A\alpha}{S_A\alpha + x} - \frac{S_B\alpha}{S_B\alpha + x},$$

so that $\Delta\Pi(d) = h(d^{\sigma/2})$. It suffices to characterise h on $[0, \infty)$. Clearly $h(0) = 1 - 1 = 0$ and $\lim_{x\uparrow\infty} h(x) = 0$. For $x > 0$, $h(x) > 0$ iff $S_A\alpha(S_B\alpha + x) > S_B\alpha(S_A\alpha + x)$, which simplifies to $(S_A - S_B)x > 0$ and holds under $S_A > S_B$.

Differentiating,

$$h'(x) \;=\; -\frac{S_A\alpha}{(S_A\alpha + x)^2} + \frac{S_B\alpha}{(S_B\alpha + x)^2}.$$

Setting $h'(x) = 0$ gives $S_A(S_B\alpha + x)^2 = S_B(S_A\alpha + x)^2$. Taking positive square roots (both sides strictly positive for $x \geq 0$),

$$\sqrt{S_A}(S_B\alpha + x) \;=\; \sqrt{S_B}(S_A\alpha + x),$$

which rearranges to $(\sqrt{S_A} - \sqrt{S_B})\, x = \alpha\sqrt{S_A S_B}(\sqrt{S_A} - \sqrt{S_B})$. Since $\sqrt{S_A} > \sqrt{S_B}$, the unique interior critical point is $x^{\text{peak}} = \alpha\sqrt{S_A S_B}$. Because h is continuous on $[0, \infty)$, vanishes at both endpoints, and is strictly positive on the interior, the unique critical point must be a strict global maximum, so h is single-peaked. Translating back via $x = d^{\sigma/2}$ yields Equation 20.

At x^{peak},

$$\frac{S_A\alpha}{S_A\alpha + \alpha\sqrt{S_A S_B}} \;=\; \frac{S_A}{S_A + \sqrt{S_A S_B}} \;=\; \frac{\sqrt{S_A}}{\sqrt{S_A} + \sqrt{S_B}},$$

and similarly $S_B\alpha/(S_B\alpha + \alpha\sqrt{S_A S_B}) = \sqrt{S_B}/(\sqrt{S_A} + \sqrt{S_B})$. Subtracting yields Equation 21. The expression depends only on the ratio S_A/S_B: dividing numerator and denominator by $\sqrt{S_B}$ gives $(\sqrt{S_A/S_B} - 1)/(\sqrt{S_A/S_B} + 1)$. ■

Lemma 3 gives a sharp bound on the Hawk's contest-based electoral advantage. For $d < d^{\text{peak}}$, every additional unit of symbolic policy amplifies the Hawk's relative advantage

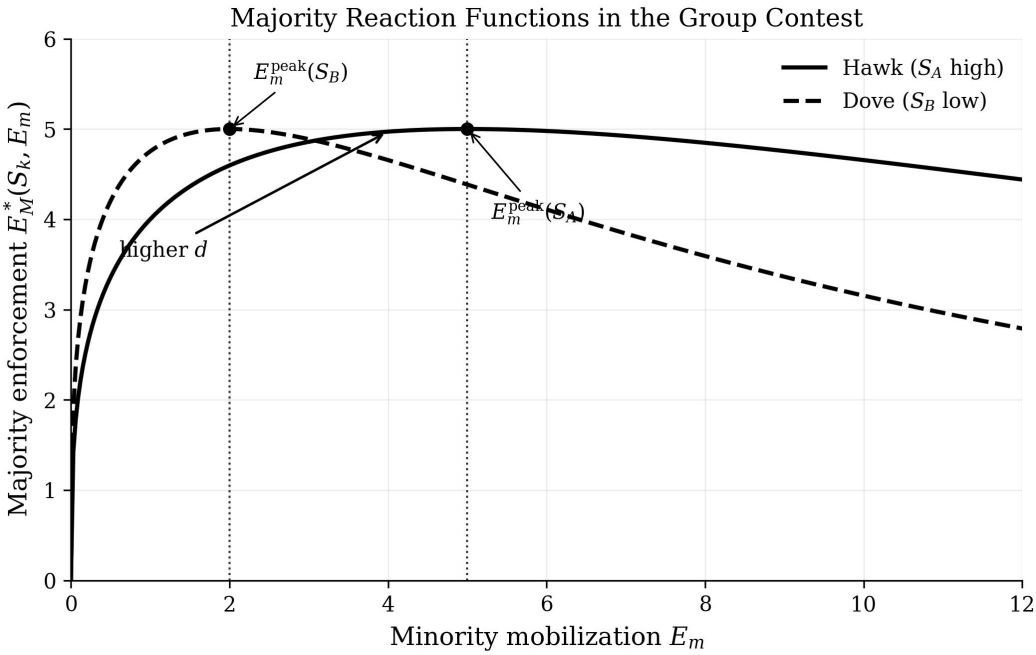

FIGURE 2 Security Premium and the provocation ceiling.

Notes: The figure plots the Security Premium $\Delta\Pi(d) = \Pi_A^*(d, S_A) - \Pi_B^*(d, S_B)$ from Equation 19 as a function of symbolic policy intensity d, for a Hawk-Dove pair with $S_A > S_B$. The curve is zero at $d = 0$ (dormant contest: both parties deliver full dominance), rises on $(0, d^{\text{peak}})$ as symbolic policy activates the contest and exposes the Hawk's technological advantage, and falls on $(d^{\text{peak}}, \infty)$ as further provocation erodes realized security for both parties. The peak value $(\sqrt{S_A} - \sqrt{S_B})/(\sqrt{S_A} + \sqrt{S_B})$ depends only on the technology ratio S_A/S_B. Illustration uses $(N_M, N_m, \Lambda, \kappa, \phi, \sigma, S_A, S_B) = (100, 100, 1, 1, 1, 1, 1.0, 0.4)$.

through the contest channel, because more intense Minority mobilization crowds in more Majority mobilization and exposes the higher productivity of the Hawk's state capacity. For $d > d^{\text{peak}}$, additional provocation is self-defeating: it drives both parties' equilibrium security toward zero faster than it widens the gap, so the Hawk's contest advantage shrinks. Even a party that seeks to manufacture conflict must calibrate provocation below this ceiling. The ceiling value $(\sqrt{S_A} - \sqrt{S_B})/(\sqrt{S_A} + \sqrt{S_B})$ is the sharpest statement of the Hawk's contest-based advantage: even if the Hawk could choose d freely and costlessly, the electoral return from symbolic policy is bounded by the pure ratio of technological capacities and is structurally independent of how grievance-elastic the Minority happens to be, how large the two populations are, or how the Majority's collective action problem scales.

2.4.3 The inefficiency of the contest: rent dissipation under biased design

Contest theory often measures the resource cost of conflict by the sum of equilibrium efforts (rent dissipation). In my environment, the incumbent's symbolic policy is the primitive "design" instrument that determines the contest intensity by shifting the joint equilibrium $(E_M^*(d, S_k), E_m^*(d, S_k))$.

Definition 1 (Rent dissipation) *For an incumbent with technology S_k, define total rent dissipation under symbolic policy d as the sum of the two sides' equilibrium efforts in the contest subgame,*

$$RD(d, S_k) \;\equiv\; E_M^*(d, S_k) + E_m^*(d, S_k). \tag{22}$$

Using Lemma 1 and $E_M^*/E_m^* = \alpha d^{-\sigma/2}$, rent dissipation admits the closed form

$$RD(d, S_k) \;=\; \left(1 + \alpha d^{-\sigma/2}\right) E_m^*(d, S_k) \;=\; \beta^2 \, \frac{S_k^{1/2} \, d^{3\sigma/4} + \alpha S_k^{1/2} \, d^{\sigma/4}}{S_k \alpha + d^{\sigma/2}}, \tag{23}$$

where $\beta \equiv (\kappa \phi N_m \Lambda N_M)^{1/4}$ as before.

Proposition 1 (Rent dissipation and biased contest design) *For any $S_k > 0$, rent dissipation satisfies $RD(0, S_k) = 0$ and $RD(d, S_k) > 0$ for all $d > 0$. Moreover, the politically optimal symbolic policy generically features strictly positive dissipation: in any equilibrium in which the incumbent chooses $d_k^* > 0$, the realized contest dissipates resources $RD(d_k^*, S_k) > 0$. Whenever symbolic policy induces a non-trivial threat, equilibrium mobilization generates rent dissipation through both opposition effort and enforcement.*

Proof. Lemma 1 establishes that $E_M^*(0, S_k) = E_m^*(0, S_k) = 0$ (dormant case), so $RD(0, S_k) = 0$. For $d > 0$, Equation 12 gives $E_m^*(d, S_k) > 0$ and $E_M^*(d, S_k) > 0$, hence $RD(d, S_k) > 0$. The final statement follows immediately: if an incumbent chooses $d_k^* > 0$ in equilibrium, then the induced dissipation is $RD(d_k^*, S_k) > 0$ by the preceding argument. In particular, combining this observation with Proposition Proposition 2 implies that strategic provocation ("arming the enemy") in equilibrium is necessarily accompanied by strictly positive rent dissipation. ■

Proposition Proposition 1 records that any equilibrium with $d_k^* > 0$ necessarily features positive rent dissipation, because d triggers both sides of the contest to exert strictly positive effort. In this sense, symbolic policy acts as a biased contest-design instrument that shifts political competition toward security at the cost of real resource waste.

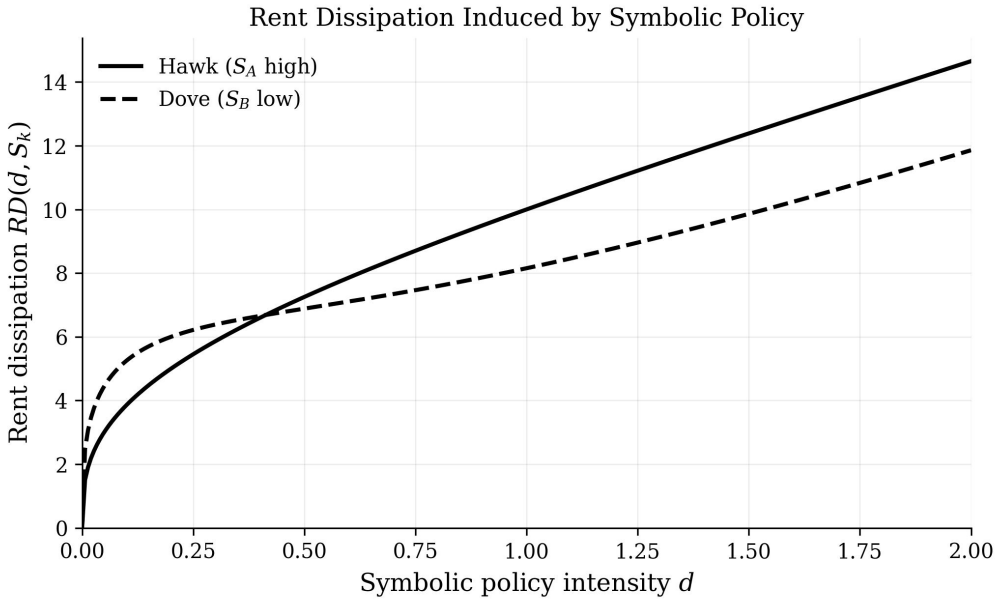

FIGURE 3 Rent dissipation under symbolic policy.

Notes: The figure plots $RD(d, S_k)$ from Equation 23 as a function of symbolic policy intensity d, for two levels of state capacity. Symbolic policy activates the joint contest in Lemma 1, generating real resource dissipation through both Minority mobilization and Majority enforcement. Unlike the neutral contest designer in standard contest models (e.g., Barut and Kovenock (1998)), the political incumbent may optimally choose a strictly positive d, and hence positive dissipation, because conflict increases the Security Premium in elections. Parameters are $(N_M, N_m, \Lambda, \kappa, \phi, \sigma, S_A, S_B) = (100, 100, 1, 1, 1, 1, 1.0, 0.2)$.

Taken together, by choosing d_k the incumbent activates the joint contest in Lemma 1 and thereby manipulates the equilibrium mapping from security competence S_k into realized security $\Pi_k^*(d_k, S_k)$. This mechanism is central to the political economy analysis below.

3 Voter Preferences and Electoral Competition

Preferences in the Majority electorate are characterized by heterogeneity. Voter i derives utility from a weighted combination of physical security and economic consumption, augmented by an idiosyncratic partisan bias. Let d denote the symbolic-policy environment currently in place: it is fixed by the incumbent at Stage 1 and pins down the joint contest equilibrium that the electorate observes at Stage 3. Within the current election, the Minority's accumulated grievance and the resulting equilibrium mobilization are sticky in d, so voters evaluate each party's security competence at the incumbent's current d rather than at

the policy the challenger would hypothetically choose upon election.[2]

The per-period utility from electing Party $k \in \{A, B\}$ under incumbency-induced environment d is

$$U_i^k(d) \;=\; \lambda\,\Pi_k^*(d, S_k) \,+\, (1-\lambda)\,u(g_k) \,-\, \chi_{\text{econ}}(d) \,-\, \chi_{\text{mod}}(d)\,\mathbb{1}[k = \text{incumbent}] \,+\, \mathbb{1}[k = A]\,\xi_i, \tag{24}$$

where $\lambda \in (0, 1)$ is a fixed (exogenous) weight on security. I decompose the cost of symbolic policy into two economically distinct channels. The *economic-disruption cost* $\chi_{\text{econ}}(d) \geq 0$ captures the real resource cost of cultural or identity-based regulations (lost productivity from ethnic segmentation, trade disruption, misallocation of enforcement capacity) and is borne by every Majority citizen regardless of who wins the next election, because the present environment is what realizes the output cost. The *moderate-defection cost* $\chi_{\text{mod}}(d) \geq 0$ captures backlash from moderate Majority voters, lost coalition support, and international-reputation costs; it is borne only while the incumbent continues the policy, so it attaches to the incumbent party in the voter's counterfactual comparison. Both functions are continuously differentiable, strictly increasing, strictly convex in d, and satisfy $\chi_{\text{econ}}(0) = \chi_{\text{mod}}(0) = 0$. I write the combined cost $\chi(d) \equiv \chi_{\text{econ}}(d) + \chi_{\text{mod}}(d)$ whenever the decomposition is not needed. The term ξ_i captures an idiosyncratic preference shock in favor of the Hawk, distributed according to a cumulative distribution function $\Phi(\cdot)$ with median zero (a normalization).

A citizen i votes for Party A if $U_i^A(d) \geq U_i^B(d)$. In the baseline scenario of interest the Hawk is the incumbent (so $k = A$ bears $\chi_{\text{mod}}(d)$), and both parties are evaluated at the common environment d. Substituting (24), the indifference condition for the *swing voter* $\hat{\xi}$ is:

$$\hat{\xi} \;=\; \underbrace{(1-\lambda)\big[u(g_B) - u(g_A)\big]}_{\text{Economic deficit}} \,+\, \underbrace{\chi_{\text{mod}}(d)}_{\text{Moderate-defection cost}} \,-\, \underbrace{\lambda\,\Delta\Pi(d)}_{\text{Security premium}}\;, \tag{25}$$

where $\Delta\Pi(d) \equiv \Pi_A^*(d, S_A) - \Pi_B^*(d, S_B)$ is the common-d Security Premium of Lemma 3. The economic-disruption cost $\chi_{\text{econ}}(d)$ is borne under either continuation and therefore cancels in the comparison. The first term captures the Hawk's economic deficit; the second is the direct administrative backlash that the Majority associates with the incumbent Hawk's provocation; the third is the Security Premium. Party A secures an electoral victory if the

2. This is the standard retrospective-voting reading in models with policy inertia: the incumbent's provocation creates a mobilization environment that cannot be unwound within the election horizon, so the voter compares each party's competence in handling *the existing environment*. Under the alternative "own-policy" reading, the Dove's dominant strategy $d_B = 0$ (Lemma 4 below) would make the Dove deliver full dominance $\Pi_B^*(0, S_B) = 1$ and the Hawk's provocation incentive would trivially collapse. The common-d reading is therefore both substantively realistic and necessary for strategic provocation to be a meaningful political phenomenon.

median voter prefers A, corresponding to $\Phi(-\hat{\xi}) > 1/2$.

4 Baseline Analysis

The political equilibrium is the subgame perfect equilibrium of the three-stage game described above. In the policy stage, the incumbent party chooses the symbolic-policy environment $d \geq 0$, taking (S_k, g_k) as given party-type characteristics. In the contest stage, joint mobilization $(E_M^*(d, S_k), E_m^*(d, S_k))$ and equilibrium security $\Pi_k^*(d, S_k)$ are determined as the symmetric Nash equilibrium of Lemma 1. In the election stage, the Majority electorate votes on the basis of indirect utilities at the current environment d: the common-d swing-voter threshold $\hat{\xi}$ in Equation 25 pins down the re-election probability, and the Security Premium that determines it is $\Delta\Pi(d)$ of Lemma 3. I solve the game by backward induction, first establishing the Dove's continuation behaviour and then characterizing the Hawk's provocation incentive.

The Dove's dominant strategy

Before turning to the Hawk's provocation problem, I first verify that the claim "the Dove does not provoke" is an *equilibrium outcome* of the Dove's own optimization, not an assumption imposed on the model. The key insight is that under the common-d reading, a positive d simultaneously (i) raises the Security Premium $\Delta\Pi(d)$ — which benefits only the Hawk, since $S_A > S_B$ — and (ii) imposes the moderate-defection cost $\chi_{\text{mod}}(d)$ on the incumbent. Neither force works in the Dove's favour, so the Dove's optimal provocation is the corner.

Lemma 4 (Dove's dominant strategy) *Suppose the Dove is the incumbent and chooses $d_B \geq 0$ to maximize its re-election probability. Under Equation 24 and the common-d swing threshold, the Dove's unique best response is $d_B^* = 0$. Consequently, when the Dove is incumbent, the joint contest is dormant in equilibrium, $(E_M^*, E_m^*) = (0, 0)$, and realized security converges to $\Pi_B^*(0, S_B) = 1$. When a Dove is in office, symbolic provocation is strictly dominated: the Dove would at best waste its higher S_B-relative-to-S_A disadvantage and bear a positive moderate-defection cost.*

Proof. Suppose the Dove is the incumbent. Retracing the derivation of Equation 25 with the moderate-defection cost now attached to $k = B$, the swing-voter indifference condition — the threshold in ξ_i above which the voter prefers the Hawk over the Dove — becomes

$$\hat{\xi}^B(d) \;=\; (1-\lambda)\,\Delta u \;-\; \chi_{\text{mod}}(d) \;-\; \lambda\,\Delta\Pi(d) \;=\; (1-\lambda)\,\Delta u \;-\; G(d),$$

where $G(d) \equiv \chi_{\text{mod}}(d) + \lambda\,\Delta\Pi(d)$. By symmetry with the Hawk-incumbent case, the Dove's vote share is $\mathbb{P}\Big[\xi_i < \hat{\xi}^B(d)\Big] = \Phi\big(\hat{\xi}^B(d)\big)$, strictly increasing in $\hat{\xi}^B(d)$. The Dove therefore maximizes its vote share by maximizing $\hat{\xi}^B(d)$, and since $(1-\lambda)\Delta u$ is independent of d, equivalently by *minimizing* $G(d)$.

By Lemma 3, $\Delta\Pi(d) \geq 0$ for all $d \geq 0$ with $\Delta\Pi(0) = 0$ and $\Delta\Pi(d) > 0$ for all $d > 0$; by the assumptions on χ_{mod}, $\chi_{\text{mod}}(0) = 0$ and $\chi_{\text{mod}}(d) > 0$ for all $d > 0$. Hence for every $d > 0$,

$$G(d) \;=\; \chi_{\text{mod}}(d) + \lambda\,\Delta\Pi(d) \;>\; 0 \;=\; \chi_{\text{mod}}(0) + \lambda\,\Delta\Pi(0) \;=\; G(0).$$

Therefore $d = 0$ is the unique global minimizer of G on $[0, \infty)$, and $d_B^* = 0$ is the Dove's unique best response. At $d_B^* = 0$, Lemma 1 yields $(E_M^*, E_m^*) = (0, 0)$ and $\lim_{d \downarrow 0} \Pi_B^*(d, S_B) = 1$. ■

Lemma 4 has two interpretive consequences. First, the Dove's corner is an endogenous equilibrium outcome, not a modelling shortcut: the Dove has no stake in activating the contest because the underlying state-capacity technology is structurally biased in the Hawk's favor. Second, the "sticky-environment" reading of Equation 25 is the only one that makes the analysis non-trivial: when the Dove is in office and $d = 0$, the voter's counterfactual is the Hawk at the *same* $d = 0$, which by Lemma 1 delivers $\Pi_A^*(0, S_A) = 1$ as well, so both parties look identical on the security dimension and the purely economic comparison selects the Dove by Equation 25 whenever $\Delta u > 0$. The Dove wins peacetime elections not because its security competence is high in absolute terms, but because the political environment it sets makes the security gap electorally invisible.

Strategic provocation by the Hawk

I now consider the Hawk's incumbency. Under Lemma 4, the relevant off-path threat is a Dove at $d = 0$; the on-path Hawk chooses $d \geq 0$ to maximize its re-election probability against this dormant benchmark. The central object is the *net electoral return* to provocation at the common environment d:

$$F(d) \;\equiv\; \lambda\,\Delta\Pi(d) \;-\; \chi_{\text{mod}}(d), \tag{26}$$

which captures the benefit of activating the security gap at the current environment, net of the moderate-defection cost the incumbent bears for doing so. The next proposition characterizes when the Hawk can survive an economic deficit and how large a deficit this channel can offset.

Proposition 2 (Strategic provocation and an economic-deficit cutoff) *Let the Hawk's economic disadvantage Δu be defined as in Equation 3. Assume χ_{mod} is continuously differentiable on $[0, \infty)$ with $\chi_{\text{mod}}(0) = 0$, $\chi'_{\text{mod}} > 0$, $\chi''_{\text{mod}} > 0$, and $\lim_{d\uparrow\infty} \chi_{\text{mod}}(d) = \infty$. Let $F(d)$ be defined as in Equation 26 and define the* provocation cutoff

$$\overline{\Delta u} \equiv \frac{1}{1-\lambda} \max_{d\geq 0} F(d). \tag{27}$$

Then:

(i) $F(0) = 0$ and F attains its maximum on $[0, \infty)$ at a finite maximizer $d_A^* \geq 0$. In particular, $\overline{\Delta u} \geq 0$.

(ii) A Hawk victory is feasible if and only if $\Delta u \leq \overline{\Delta u}$. If $\Delta u > 0$ and the Hawk wins, any optimal policy satisfies $d_A^* > 0$, and hence $E_M^*(d_A^*, S_A), E_m^*(d_A^*, S_A) > 0$: strategic provocation activates the joint contest.

(iii) The cutoff is bounded above by the pure technology ratio

$$\overline{\Delta u} \leq \frac{\lambda}{1-\lambda} \cdot \frac{\sqrt{S_A} - \sqrt{S_B}}{\sqrt{S_A} + \sqrt{S_B}}, \tag{28}$$

with equality in the limit $\chi_{\text{mod}} \to 0$.

In words: to win in spite of an economic disadvantage, the Hawk must turn on the joint contest by choosing strictly positive symbolic policy. Doing so opens a wedge between the two parties' equilibrium security at the common environment, and this wedge can partially offset the economic deficit. The maximum deficit the Hawk can overcome through provocation is finite, and its sharp bound depends only on the relative technology ratio S_A/S_B — not on the Minority's grievance elasticity, the two groups' sizes, or the convexity of cost.

Proof. *Step 1 (Existence, $F(0) = 0$, and upper bound).* By Lemma 3, $\Delta\Pi(d)$ is continuous on $[0, \infty)$ with $\Delta\Pi(0) = 0$ and bounded above by the ceiling $\overline{\Delta\Pi} \equiv (\sqrt{S_A} - \sqrt{S_B})/(\sqrt{S_A} + \sqrt{S_B})$. Hence $F(d) = \lambda\Delta\Pi(d) - \chi_{\text{mod}}(d)$ is continuous on $[0, \infty)$ with $F(0) = \lambda \cdot 0 - 0 = 0$ and, since $\chi_{\text{mod}}(d) \uparrow \infty$, $\lim_{d\uparrow\infty} F(d) = -\infty$. Therefore F attains a finite maximum $F^* \equiv \max_{d\geq 0} F(d) \in [0, \lambda\overline{\Delta\Pi}]$ at some finite $d_A^* \geq 0$. In particular $F^* \geq F(0) = 0$, so $\overline{\Delta u} \geq 0$, which establishes (i).

Step 2 (Characterization of the Hawk's optimal policy). Since the Hawk is the incumbent, Equation 24 attaches $\chi_{\text{mod}}(d)$ to the Hawk, while the Dove at $d_B^* = 0$ (from Lemma 4) bears

no moderate-defection cost. The common-d swing-voter threshold Equation 25 therefore specializes to

$$\hat{\xi}(d) \;=\; (1-\lambda)\Delta u \;+\; \chi_{\text{mod}}(d) \;-\; \lambda\,\Delta\Pi(d) \;=\; (1-\lambda)\Delta u \;-\; F(d). \tag{29}$$

The Hawk's vote share is $\Phi(-\hat{\xi}(d))$ with Φ strictly increasing, so maximizing the vote share is equivalent to minimizing $\hat{\xi}(d)$, and hence to maximizing $F(d)$ (since $(1-\lambda)\Delta u$ is independent of d). By Step 1, a finite maximizer $d_A^* \geq 0$ exists.

Step 3 (Cutoff characterization). The Hawk wins whenever the median voter (weakly) prefers A, i.e., $\hat{\xi}(d_A^*) \leq 0$. Using Equation 29, this is equivalent to $F(d_A^*) \geq (1-\lambda)\Delta u$. Since $F(d_A^*) = F^*$ by Step 2, the Hawk's re-election is feasible iff $F^* \geq (1-\lambda)\Delta u$, equivalently

$$\Delta u \;\leq\; \frac{F^*}{1-\lambda} \;=\; \overline{\Delta u},$$

which establishes the first part of (ii). If $\Delta u > 0$ and the Hawk wins, then $F(d_A^*) \geq (1-\lambda)\Delta u > 0$; combined with $F(0) = 0$, this rules out $d_A^* = 0$, so $d_A^* > 0$. By Lemma 1, $d_A^* > 0$ implies $E_m^*(d_A^*, S_A), E_M^*(d_A^*, S_A) > 0$ — the joint contest is active along the strategic-provocation path.

Step 4 (Technology-ratio ceiling). From $F(d) = \lambda\Delta\Pi(d) - \chi_{\text{mod}}(d) \leq \lambda\overline{\Delta\Pi}$ for all d, one gets $F^* \leq \lambda\overline{\Delta\Pi}$, and dividing by $(1-\lambda)$ yields Equation 28. For the equality-in-the-limit statement, take any vanishing sequence of cost functions $\chi_{\text{mod}}^{(n)} \to 0$ pointwise with bounded variation. Then $F^{(n)}(d) \to \lambda\Delta\Pi(d)$ pointwise, and in particular $F^{(n)}(d^{\text{peak}}) \to \lambda\overline{\Delta\Pi}$ where $d^{\text{peak}} = (\alpha\sqrt{S_A S_B})^{2/\sigma}$ is the Security-Premium maximizer from Equation 20. Hence $F^{*,(n)} \to \lambda\overline{\Delta\Pi}$ and $\overline{\Delta u}^{(n)} \to \frac{\lambda}{1-\lambda}\overline{\Delta\Pi}$, which is the right-hand side of Equation 28. ■

Remark 1 (Provocation ceiling) *Equation 28 is the sharpest statement of the Hawk's contest-based advantage: regardless of the convexity of costs, the two groups' sizes, or how grievance-elastic the Minority is, the maximum economic deficit the Hawk can offset through symbolic provocation is a monotone function only of the technology ratio S_A/S_B. In the costless-provocation limit $\chi_{\text{mod}} \to 0$, this bound binds at the Security Premium maximizer d^{peak}, so the Hawk's optimal provocation converges to the contest-theoretic "peak" of Lemma 3.*

Proposition 2 shows identity politics as a substitute for competence and its limit. The Hawk is structurally disadvantaged on economic performance, so it cannot win by competing on the public-good dimension. In a peaceful environment, this disadvantage is critical in elections because the underlying contest is dormant: when $d = 0$, Lemma 1 gives

$\Pi_A^*(0, S_A) = \Pi_B^*(0, S_B) = 1$, so the Hawk's state-capacity advantage is electorally invisible and the Hawk loses the election once their economic incompetence is revealed. To survive, the Hawk starts provocation. By choosing $d_A^* > 0$, the Hawk activates the joint contest and triggers mobilization on both sides, so the equilibrium security of each technology becomes electorally relevant at the common environment. The Hawk's provocation makes the security dimension salient in the swing voter's calculus, and its higher S_A then translates into a positive Security Premium $\Delta\Pi(d_A^*) > 0$.

Term $\overline{\Delta u}$ captures how much voters will sacrifice in economic payoff for identity security. When $\Delta u > \overline{\Delta u}$, the electorate's tolerance for economic loss breaks. Restoring consumption dominates any additional security advantage the Hawk can create with identity contestation. Under these mechanisms, a significant negative economic shock can moderate identity politics in an electoral economy. 1 sharpens this intuition: the maximum economic deficit the Hawk can offset is itself a monotone function only of the technology ratio S_A/S_B, so societies in which the two parties differ strongly in security competence permit a wider wedge of tolerable economic failure.

5 Dynamic Extension: Myopia, Forgetting, and Regime Switching

This section develops a parsimonious dynamic extension that delivers political-economy cycles, alternating high-d/low-g and low-d/high-g regimes, without allowing the preference weight on security λ to vary over time. The key ingredients are (i) aggregate shocks that are imperfectly observed, (ii) retrospective (myopic) electoral accountability, and (iii) the regime-dependence of conflict production through S_k and the induced equilibrium mobilization $E_M^*(S_k, E_m)$.

5.1 Aggregate shocks and noisy performance

Time is discrete, $t = 0, 1, 2, \ldots$, and an election takes place each period. Let θ_t denote an aggregate economic state and ω_t an aggregate threat state. For expositional simplicity, take party characteristics (S_k, g_k, d_k) as stationary over time, with $g_B > g_A$ and $d_A > d_B$.

Under a candidate party $k \in \{A, B\}$, the realized economic payoff is

$$c_{k,t} = \theta_t g_k, \tag{30}$$

and minority mobilization is

$$E_{m,t}(k) = \omega_t \gamma d_k^\sigma. \tag{31}$$

Given $(S_k, E_{m,t}(k))$, Majority mobilization is $E^*_{M,t}(S_k, E_{m,t}(k))$, and security is $\Pi_{k,t} = \Pi_k(E^*_{M,t}, E_{m,t}(k))$ as in Equation 4.

Voters do not observe (θ_t, ω_t) nor the counterfactual outcomes under the challenger. Instead, they observe noisy performance signals of consumption and disorder under the incumbent regime k_t:

$$\tilde{c}_t = c_{k_t,t} + \varepsilon^c_t, \qquad \tilde{v}_t = \big(1 - \Pi_{k_t,t}\big) + \varepsilon^v_t, \tag{32}$$

where $\tilde{v}_t$ is a reduced-form measure of observable street disorder/violence. Importantly, because $\Pi_{k,t}$ depends on S_k and the endogenous response $E^*_{M,t}(S_k, E_m)$, the same underlying threat ω_t can generate very different observed disorder across regimes.

5.2 Forward-looking voting and the conflict trap

Depending on voter patience, the dynamic environment admits two qualitatively distinct equilibrium paths. When voters are forward-looking, regime choice is governed by the discounted value of future security capacity.

Proposition 3 (Conflict Trap) *Assume voters are forward-looking with discount factor $\delta \in [0, 1)$ and form rational expectations about the future evolution of (θ_t, ω_t) under each regime. There exists a critical patience threshold $\bar{\delta} \in (0, 1)$ such that for all $\delta > \bar{\delta}$, the Hawk regime is an absorbing state in a stationary political equilibrium. In particular, conditional on a Hawk incumbent, the electorate retains the Hawk even after the most adverse economic realizations $\theta_t = \underline{\theta}$ in order to avoid the discrete reduction in security capacity associated with switching from S_A to S_B. Sufficiently patient voters tolerate even severe short-run economic losses to avoid a persistent drop in future security capacity.*

Proof. Fix an incumbent Hawk regime and suppress the idiosyncratic taste term ξ_i by considering the median voter (normalized to have $\xi_i = 0$). Let

$$\Omega(\theta, \omega) \;\equiv\; \lambda\Big(\Pi^*_A(S_A, \omega\gamma d^\sigma_A) - \Pi^*_B(S_B, \omega\gamma d^\sigma_B)\Big) - (1-\lambda)\Big(u(\theta g_B) - u(\theta g_A)\Big) - \big[\chi(d_A) - \chi(d_B)\big] \tag{33}$$

denote the one-period net advantage of the Hawk relative to the Dove in state (θ, ω). Consider the worst economic realization $\underline{\theta}$ and assume (as in the proposition) that it constitutes the most unfavorable state for the Hawk's re-election incentive.

To isolate the role of patience, consider a transitory recession at date t with $\theta_t = \underline{\theta}$, followed by reversion from $t + 1$ onward to a stationary environment. In that continuation

environment, the discrete gap in security capacity ($S_A > S_B$) implies that the expected per-period net advantage of the Hawk is strictly positive:

$$\overline{\Omega} \;\equiv\; \mathbb{E}\big[\Omega(\theta,\omega) \mid t+1 \text{ onward}\big] \;>\; 0. \tag{34}$$

Let $\underline{\Omega} \equiv \mathbb{E}\big[\Omega(\underline{\theta},\omega_t)\big]$ denote the expected one-period net advantage during the recession. The forward-looking value difference between retaining the Hawk and switching regimes at date t can then be expressed (via the Bellman equations under stationarity from $t+1$ onward) as

$$\Delta V_t \;=\; \underline{\Omega} + \sum_{\tau=t+1}^{\infty} \delta^{\tau-t}\overline{\Omega} \;=\; \underline{\Omega} + \frac{\delta}{1-\delta}\overline{\Omega}. \tag{35}$$

The function $\Delta V_t(\delta)$ is continuous and strictly increasing on $[0,1)$. Moreover, $\lim_{\delta\uparrow 1} \Delta V_t(\delta) = +\infty$ by Equation 34. Hence there exists a cutoff $\bar{\delta} \in (0,1)$ such that $\Delta V_t(\delta) \geq 0$ for all $\delta \geq \bar{\delta}$. For $\delta > \bar{\delta}$ the median voter weakly prefers to retain the Hawk even when $\theta_t = \underline{\theta}$. Since $\underline{\theta}$ is the most adverse economic realization, the Hawk is retained in every state, implying that the Hawk regime is absorbing in a stationary equilibrium. ■

Proposition 3 demonstrates how patience of voters interacts with accountability. Voters internalize regime switching from Hawk to Dove put them into weaker technologies of identity protection. A weak S entails weaker ability of the majority to coordinate under future threats of minority. When δ is sufficient high, the discounted value of avioding future risks dominates even the worst contemporaneous economi shock $\underline{\theta}$. Current electorial punishment for economic underperformance of the Hawk fades because of voter's future concerns over identity security. This mechanism captures why identity-fragmented societies be locked into a conflict trap in which the Hawk survive indefinitely long despite worse economic management.

5.3 Myopic voting and regime transitions

Let per-period utility be given by (24). A forward-looking voter with discount factor $\delta \in [0,1)$ evaluates parties by lifetime value. The special case $\delta \to 0$ corresponds to purely retrospective (myopic) voting: the electorate selects the party that maximizes perceived one-period utility based on the current performance signals $(\tilde{c}_t, \tilde{v}_t)$.

To connect the dynamic environment to the static swing-voter logic, define the perceived

one-period Hawk-selection condition

$$\lambda\Big(\mathbb{E}_t[\Pi_{A,t}] - \mathbb{E}_t[\Pi_{B,t}]\Big) \;\geq\; (1-\lambda)\Big(\mathbb{E}_t[u(\theta_t g_B)] - \mathbb{E}_t[u(\theta_t g_A)]\Big) + \big[\chi(d_A) - \chi(d_B)\big], \tag{36}$$

which is the dynamic analogue of Equation 25.

Proposition 4 (Political Cycles) *Assume voters are myopic ($\delta = 0$), and (θ_t, ω_t) are stochastic with full support on $[\underline{\theta}, \overline{\theta}] \times [\underline{\omega}, \overline{\omega}]$. Suppose (i) the posterior expectations entering Equation 36 are monotone in the underlying states,*[3] *(ii) for $g_B > g_A$, the function $\theta \mapsto u(\theta g_B) - u(\theta g_A)$ is strictly decreasing on $[\underline{\theta}, \overline{\theta}]$, and (iii) the Security Premium is increasing in minority mobilization over the policy-relevant domain (Lemma Lemma 3). Assume further that the myopic election rule selects different parties in different states (i.e., there exist (θ, ω) and (θ', ω') in the support such that Equation 36 holds in one and fails in the other). Then the political equilibrium exhibits recurrent regime switching. In particular:*

(i) (Hawk $\to$ Dove) *Conditional on a Hawk incumbent, the probability of a regime change is strictly decreasing in θ_t.*

(ii) (Dove $\to$ Hawk) *Conditional on a Dove incumbent, the probability of a regime change is strictly increasing in the threat environment ω_t, the minority's structural mobilization potential γ, and the security weight λ.*

Myopic retrospective voting makes regime choice track current conditions: adverse economics push the electorate toward the Dove, while high perceived threat pushes it toward the Hawk.

Proof. Under myopic voting ($\delta = 0$), the median voter elects the Hawk at date t if and only if Equation 36 holds. Define the net advantage

$$\Omega_t \;\equiv\; \lambda\Big(\mathbb{E}_t[\Pi_{A,t}] - \mathbb{E}_t[\Pi_{B,t}]\Big) - (1-\lambda)\Big(\mathbb{E}_t[u(\theta_t g_B)] - \mathbb{E}_t[u(\theta_t g_A)]\Big) - \big[\chi(d_A) - \chi(d_B)\big], \tag{37}$$

so the Hawk is elected if and only if $\Omega_t \geq 0$.

For part (i), note that for $g_B > g_A$,

$$\frac{\partial}{\partial\theta}\Big(u(\theta g_B) - u(\theta g_A)\Big) = g_B u'(\theta g_B) - g_A u'(\theta g_A), \tag{38}$$

3. For instance, a natural sufficient condition is that the signal structure $(\tilde{c}_t, \tilde{v}_t)$ satisfies a monotone likelihood ratio property so that $\mathbb{E}_t[\theta_t]$ and $\mathbb{E}_t[\omega_t]$ are monotone in the realized signals.

so assumption (ii) is equivalent to $g_B u'(\theta g_B) - g_A u'(\theta g_A) < 0$ on $[\underline{\theta}, \overline{\theta}]$.[4] Combining this with assumption (i), a lower realization of θ_t strictly increases the perceived economic advantage of the Dove and therefore strictly decreases Ω_t.

Moreover, Ω_t is continuous in (θ_t, ω_t) (given continuity of u and monotone posteriors), and by hypothesis there exist states in which $\Omega_t \geq 0$ and states in which $\Omega_t < 0$. Fix any $\omega \in [\underline{\omega}, \overline{\omega}]$ and hold $\omega_t = \omega$. The previous argument implies that Ω_t is strictly increasing in θ_t on $[\underline{\theta}, \overline{\theta}]$. Hence the intermediate value theorem yields a cutoff $\theta^*(\omega) \in (\underline{\theta}, \overline{\theta})$ such that the myopic electorate selects the Hawk if and only if $\theta_t \geq \theta^*(\omega_t)$. Therefore, conditional on a Hawk incumbent, the Hawk $\to$ Dove transition event $\{\Omega_t < 0\}$ is equivalent to $\{\theta_t < \theta^*(\omega_t)\}$, which becomes (weakly) more likely under more adverse economic realizations.

For part (ii), note that minority mobilization satisfies $E_{m,t}(k) = \omega_t \gamma d_k^\sigma$, so higher ω_t or higher γ increases the threat faced under either candidate platform. By Lemma Lemma 3 (over the relevant domain), the perceived Security Premium $\mathbb{E}_t[\Pi_{A,t} - \Pi_{B,t}]$ is increasing in the threat level, so Ω_t is increasing in ω_t and γ. Moreover,

$$\frac{\partial \Omega_t}{\partial \lambda} = \Big(\mathbb{E}_t[\Pi_{A,t}] - \mathbb{E}_t[\Pi_{B,t}]\Big) + \Big(\mathbb{E}_t[u(\theta_t g_B)] - \mathbb{E}_t[u(\theta_t g_A)]\Big) > 0, \tag{39}$$

so Ω_t is increasing in λ. Therefore the probability that Ω_t crosses zero in favor of the Hawk (a Dove $\to$ Hawk transition) is increasing in ω_t, γ, and λ.

Finally, since the state process has full support and (by assumption) there exist states in which $\Omega_t \geq 0$ and states in which $\Omega_t < 0$, both transitions occur with positive probability along the realized path, yielding recurrent regime switching. ■

Myopic voting breaks the electoral trap by limiting voter's horizon to current signals. When a Hawk incumbent successfully suppress visible disorder, voters infer (contemporaneously) current environment is peaceful but fails to account for the peace is sustained by high security capacity by the Hawk party. Under this equilibrium, a sufficiently adverse economic realization triggers the electorate to vote for the Dove party to recover public goods provision in the economy. When the Dove is incumbent, lower capacity in identity security weakens Majority mobilization, reduces deterrence, and the safety equilibrium cannot be sustained. Identity threat from the Minority triggers a swing back to the Hawk to contain the potential disorder with cultural policies.

The resulting cycles can be read as over-correction as voters demand the dimension that is currently scarce. Economic prosperity is demanded in peaceful periods and security in

4. A common sufficient condition is CRRA utility $u(c) = c^{1-\rho}/(1-\rho)$ with $\rho > 1$, in which case $g_B u'(\theta g_B) - g_A u'(\theta g_A) = \theta^{-\rho}\big(g_B^{1-\rho} - g_A^{1-\rho}\big) < 0$ for all $\theta > 0$.

identity is demanded in turbulent periods. The economy therefore oscillates between cycles of political control due to the presence of the minority.

Corollary 2 (The Paradox of Competence) *Under myopic voting ($\delta = 0$), a first-order stochastic decrease in observed disorder under a Hawk incumbent (lower $\tilde{v}_t$) weakly decreases the perceived Security Premium in Equation 36 and thereby weakly increases the probability of a Hawk defeat. Consequently, conditional on a Hawk incumbent, the hazard of regime change is decreasing in observed disorder. A more successful Hawk can look less necessary in equilibrium: quieter streets shrink the perceived Security Premium and can increase turnover risk.*

Proof. Under myopic voting, the Hawk is re-elected at date t if and only if Equation 36 holds, equivalently if $\Omega_t \geq 0$ where Ω_t is defined in the proof of Proposition Proposition 4. By assumption (i) in Proposition Proposition 4 (and the monotone-likelihood-ratio sufficient condition in its footnote), a lower realization of the disorder signal $\tilde{v}_t$ weakly shifts posterior beliefs toward a lower threat environment and therefore weakly decreases the perceived Security Premium $\mathbb{E}_t[\Pi_{A,t} - \Pi_{B,t}]$.

Since Ω_t is increasing in $\mathbb{E}_t[\Pi_{A,t} - \Pi_{B,t}]$, a first-order stochastic decrease in $\tilde{v}_t$ induces a first-order stochastic decrease in Ω_t and hence weakly decreases $\mathbb{P}(\Omega_t \geq 0)$, the re-election probability of the Hawk. Equivalently, the probability of a regime change, $\mathbb{P}(\Omega_t < 0)$, weakly increases when observed disorder is stochastically lower. ■

Corollary Corollary 2 highlights an inference effect: under myopia, low observed disorder under the Hawk shifts beliefs toward a benign threat environment and reduces the perceived Security Premium, raising the probability of a Hawk defeat.

Together, Propositions Proposition 3 and Proposition 4 show that the same micro-foundation, namely a biased incumbent manipulating a group contest to solve a collective action problem, can generate either long-run entrenchment or recurrent turnover. Patience makes the electorate internalize the durable loss of security capacity from regime change, producing an absorbing Hawk equilibrium. Myopia makes voters treat security as a contemporaneous outcome, opening the door to regime switching driven by noisy shocks and retrospective accountability.

References

Barut, Yaşar and Dan Kovenock, "The symmetric multiple prize all-pay auction with complete information," *European Journal of Political Economy*, 1998, *14* (4), 627–644.

Fang, Lily, Thomas H. Noe, and Philipp Strack, "Turning Up the Heat: The Discouraging Effect of Competition in Contests," *Journal of Political Economy*, 2020, *128* (5), 1940–1975.

Jia, Hao, "A stochastic derivation of the ratio form of contest success functions," *Public Choice*, 2008, *135* (3), 125–130.

Skaperdas, Stergios, "Contest success functions," *Economic theory*, 1996, *7* (2), 283–290.